qSTP 2016

ASSIGNMENT BASED ON

MULTI-CLASS CLASSIFICATION AND NEURAL NETWORKS

Part B: Regular Assignment (Compulsory)

Start Date: As soon as you receive it

Due Date: Saturday June 25th, 2016 midnight

1. **ALL THE CODES HAVE TO BE IN VECTORIZATION FORMAT [EXCEPT WHERE NOT POSSIBLE]**

2. **MAKE NECESSARY ASSUMPTIONS WHEREVER REQUIRED.**

3. This is a compulsory assignment.

4. Follow all the steps very carefully!

Introduction

In this exercise, you will implement one-vs-all logistic regression and neural networks to recognize hand-written digits.

Files included in this exercise

ex3.m - Octave script that will help step you
through part 1 ex3 nn.m - Octave script that will
help step you through part 2 ex3data1.mat -
Training set of hand-written digits ex3weights.mat -
Initial weights for the neural network exercise
displayData.m - Function to help visualize the
dataset fmincg.m

- Function minimization routine (similar to fminunc) sigmoid.m

- Sigmoid function [?]

lrCostFunction.m - Logistic regression cost
function [?] oneVsAll.m - Train a one-vs-all
multi-class classifier [?]

predictOneVsAll.m - Predict using a one-vs-all multi-
class classifier [?] predict.m - Neural network
prediction function

? Indicates files you will need to complete

Throughout the exercise, you will be using the scripts ex3.m and ex3 nn.m.
These scripts set up the dataset for the problems and make calls to functions
that you will write. You do not need to modify these scripts. You are only required
to modify functions in other files, by following the instructions in this assignment.

1. **Multi-class Classification**

For this exercise, you will use logistic regression and neural networks to recognize handwritten digits (from 0 to 9). Automated handwritten digit recognition is widely used today - from recognizing zip codes (postal codes) on mail envelopes to recognizing amounts written on bank checks. This exercise will show you how the methods you've learned can be used for this classification task. In the first part of the exercise, you will extend your previous implementation of logistic regression and apply it to one-vs-all classification.

1.1 Dataset

You are given a data set in ex3data1.mat that contains 5000 training ex- amples of handwritten digits.1 The .mat format means that that the data has been saved in a native Octave/Matlab matrix format, instead of a text (ASCII) format like a csv-file. These matrices can be read directly into your program by using the load command. After loading, matrices of the correct dimensions and values will appear in your program's memory. The matrix will already be named, so you do not need to assign names to them.

```
%Load   saved   matrices
from                   file
load('ex3data1.mat');
```

```
%The matrices X and y will now be in your Octave environment
```

There are 5000 training examples in ex3data1.mat, where each training example is a 20 pixel by 20 pixel grayscale image of the digit. Each pixel is represented by a floating point number indicating the grayscale intensity at that location. The 20 by 20 grid of pixels is "unrolled" into a 400-dimensional vector. Each of these training examples become a single row in our data matrix X. This gives us a 5000 by 400 matrix X where every row is a training example for a handwritten digit image.

$$X = \begin{bmatrix} - \ (x^{(1)})^T \ - \\ - \ (x^{(2)})^T \ - \\ \vdots \\ - \ (x^{(m)})^T \ - \end{bmatrix}$$

The second part of the training set is a 5000-dimensional vector y that contains labels for the training set. To make things more compatible with Octave/Matlab indexing, where there is no zero index, we have mapped the digit zero to the value ten. Therefore, a "0" digit is labeled as "10", while the digits "1" to "9" are labeled as "1" to "9" in their natural order.

1.2 Visualizing the data

You will begin by visualizing a subset of the training set. In Part 1 of ex3.m, the code randomly selects selects 100 rows from X and passes those rows to the displayData function. This function maps each row to a 20 pixel by 20 pixel grayscale image and displays the images together. The displayData function is provided, and you are encouraged to examine the code to see how it works. After you run this step, you should see an image like:

1.3 Vectorizing Logistic Regression

You will be using multiple one-vs-all logistic regression models to build a multi-class classifier. Since there are 10 classes, you will need to train 10 separate logistic regression classifiers. To make this training efficient, it is important to ensure that your code is well vectorized. In this section, you will implement a vectorized version of logistic regression that does not employ any for loops. You can use your code in the last exercise as a starting point for this exercise.

1.3.1 Vectorizing the cost function

We will begin by writing a vectorized version of the cost function. Recall that in (unregularized) logistic regression, the cost function is:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^{m} \left[-y^{(i)} \log(h_\theta(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)})) \right]$$

To compute each element in the summation, we have to compute $h_\theta(x(i))$ for

every example i, where $h_\theta(x^{(i)}) = g(\theta^T x(i))$ and $g(z) = 1/(1+e^{-z})$ is the sigmoid

function. It turns out that we can compute this

quickly for all our examples by using matrix multiplication. Let us define X and θ

as:

$$
X = \begin{bmatrix} - (x^{(1)})^T - \\ - (x^{(2)})^T - \\ \vdots \\ - (x^{(m)})^T - \end{bmatrix} \quad \text{and} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}.
$$

Then, by computing the matrix product Xθ, we have

$$
X\theta = \begin{bmatrix} - (x^{(1)})^T\theta - \\ - (x^{(2)})^T\theta - \\ \vdots \\ - (x^{(m)})^T\theta - \end{bmatrix} = \begin{bmatrix} - \theta^T(x^{(1)}) - \\ - \theta^T(x^{(2)}) - \\ \vdots \\ - \theta^T(x^{(m)}) - \end{bmatrix}.
$$

In the last equality, we used the fact that $a^T b = b^T a$ if a and b are vectors. This

allows us to compute the products $\theta^T x^{(i)}$ for all our examples i in one line of

code.

Your job is to write the unregularized cost function in the file lrCostFunction.m
Your implementation should use the strategy we presented above to calcu-

late $\theta^T x(i)$. You should also use a vectorized approach for the rest of the cost
function. A fully vectorized version of lrCostFunction.m should not contain any
loops.

(Hint: You might want to use the element-wise multiplication operation (.*) and the
sum operation sum when writing this function)

1.3.2 Vectorizing the gradient

Recall that the gradient of the (unregularized) logistic regression cost is a vector
where the jth element is defined as:

$$\frac{\partial J}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} \right).$$

To vectorize this operation over the dataset, we start by writing out all the
partial derivatives explicitly for all θ_j,

$$
\begin{bmatrix}
\frac{\partial J}{\partial \theta_0} \\
\frac{\partial J}{\partial \theta_1} \\
\frac{\partial J}{\partial \theta_2} \\
\vdots \\
\frac{\partial J}{\partial \theta_n}
\end{bmatrix}
= \frac{1}{m}
\begin{bmatrix}
\sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x_0^{(i)} \right) \\
\sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x_1^{(i)} \right) \\
\sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x_2^{(i)} \right) \\
\vdots \\
\sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x_n^{(i)} \right)
\end{bmatrix}
$$

$$= \frac{1}{m} \sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x^{(i)} \right)$$

$$= \frac{1}{m} X^T (h_\theta(x) - y).$$

where

$$
h_\theta(x) - y =
\begin{bmatrix}
h_\theta(x^{(1)}) - y^{(1)} \\
h_\theta(x^{(2)}) - y^{(2)} \\
\vdots \\
h_\theta(x^{(1)}) - y^{(m)}
\end{bmatrix}.
$$

Note that $x^{(i)}$ is a vector, while $(h_\theta(x^{(i)}){-}y^{(i)})$ is a scalar (single number). To understand the last step of the derivation, let $\beta i = (h_\theta(x^{(i)}) - y^{(i)})$ and observe that:

$$\sum_i \beta_i x^{(i)} = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \cdots & x^{(m)} \\ | & | & & | \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_m \end{bmatrix} = X\beta ,$$

where the values $\beta_i = (h_\theta(x^{(i)}) - y^{(i)})$.

The expression above allows us to compute all the partial derivatives without any loops. If you are comfortable with linear algebra, we encourage you to work through the matrix multiplications above to convince yourself that the vectorized version does the same computations. You should now implement Equation 1 to compute the correct vectorized gradient. Once you are done, complete the function lrCostFunction.m by implementing the gradient.

Debugging Tip: Vectorizing code can sometimes be tricky. One com- mon strategy for debugging is to print out the sizes of the matrices you are working with using the size function. For example, given a data ma- trix X of size 100 × 20 (100 examples, 20 features) and θ, a vector with dimensions 20×1, you can observe that Xθ is a valid multiplication oper- ation, while θX is not. Furthermore, if you have a non-vectorized version of your code, you can compare the output of your vectorized code and non-vectorized code to make sure that they produce the same outputs.

1.3.3 Vectorizing regularized logistic regression

After you have implemented vectorization for logistic regression, you will now add regularization to the cost function. Recall that for regularized logistic regression, the cost function is defined as:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^{m} \left[-y^{(i)} \log(h_\theta(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^{n} \theta_j^2.$$

Note that you should *not* be regularizing θ_0 which is used for the bias term.

Correspondingly, the partial derivative of regularized logistic regression cost for θ_j is defined as

$$\frac{\partial J}{\partial \theta_0} = \frac{1}{m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}) \qquad \text{for } j = 0$$

$$\frac{\partial J}{\partial \theta_j} = \frac{1}{m} \left(\sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} + \lambda \theta_j \right) \qquad \text{for } j \geq 1.$$

Now modify your code in lrCostFunction to account for regularization. Once again, you should not put any loops into your code.

Octave Tip: When implementing the vectorization for regularized lo- gistic regression, you might often want to only sum and update certain elements of θ. In Octave, you can index into the matrices to access and update only certain elements. For example, A(:, 3:5) = B(:, 1:3) will replaces the columns 3 to 5 of A with the columns 1 to 3 from B. One special keyword you can use in indexing is the end keyword in indexing. This allows us to select columns (or rows) until the end of the matrix. For example, A(:, 2:end) will only return elements from the 2nd to last column of A. Thus, you could use this together with the sum and .^ op- erations to compute the sum of only the elements you are interested in (e.g., sum(z(2:end).^2)). In the starter code, lrCostFunction.m, we have also provided hints on yet another possible method computing the regularized gradient.

1.4 One-vs-all Classification

In this part of the exercise, you will implement one-vs-all classification by training multiple regularized logistic regression classifiers, one for each of the K classes in our dataset (Figure 1). In the handwritten digits dataset, K = 10, but your code should work for any value of K.

You should now complete the code in oneVsAll.m to train one classifier for each class. In particular, your code should return all the classifier parameters in a matrix $\Theta \in R^{K \times (N+1)}$, where each row of Θ corresponds to the learned logistic regression

parameters for one class. You can do this with a "for"-loop from 1 to K, training each classifier independently.

Note that the y argument to this function is a vector of labels from 1 to 10, where the digit "0" is mapped to the label 10 (to avoid confusions with indexing).

When training the classifier for class $k \in \{1,...,K\}$, you will want a m- dimensional vector of labels y, where $y_j \in 0,1$ indicates whether the j-th training instance belongs to class k ($y_j = 1$), or if it belongs to a different class ($y_j = 0$). You may find logical arrays helpful for this task.

<u>Octave Tip</u>: Logical arrays in Octave are arrays which contain binary (0 or 1) elements. In Octave, evaluating the expression a == b for a vector a (of size m×1) and scalar b will return a vector of the same size as a with ones at positions where the elements of a are equal to b and zeroes where they are different. To see how this works for yourself, try the following code in Octave:

a = 1:10; % Create
a and b b = 3;

a == b % You should try different values of b here

Furthermore, you will be using fmincg for this exercise (instead of fminunc). fmincg works similarly to fminunc, but is more more efficient for dealing with a large number of parameters.

After you have correctly completed the code for oneVsAll.m, the script ex3.m will continue to use your oneVsAll function to train a multi-class clas- sifier.

1.4.1 One-vs-all Prediction

After training your one-vs-all classifier, you can now use it to predict the digit contained in a given image. For each input, you should compute the "probability" that it belongs to each class using the trained logistic regression classifiers. Your one-vs-all prediction function will pick the class for which the corresponding logistic regression classifier outputs the highest probability and return the class label (1, 2,..., or K) as the prediction for the input example.

You should now complete the code in predictOneVsAll.m to use the one-vs-all classifier to make predictions.

Once you are done, ex3.m will call your predictOneVsAll function using the learned value of Θ. You should see that the training set accuracy is about 94.9% (i.e., it classifies 94.9% of the examples in the training set correctly).

2. __Neural Networks__

In the previous part of this exercise, you implemented multi-class re-gression to recognize handwritten digits. However, logistic regression cannot form more complex hypotheses as it is only a linear classifier.2 In this part of the exercise, you will implement a neural network to rec- ognize handwritten digits using the same training set as before. The neural network will be able to represent complex models that form non-linear hy- potheses. You will be using parameters from a neural network that we have already trained. Your goal is to implement the feedforward propagation algorithm to use our weights for prediction. In next week's ex- ercise, you will write the backpropagation algorithm for learning the neural network parameters. The provided script, ex3 nn.m, will help you step through this exercise.

2.1 Model representation

Our neural network is shown in the figure below. It has 3 layers – an input layer, a hidden layer and an output layer. Recall that our inputs are pixel values of digit images. Since the images are of size 20×20, this gives us 400 input layer units (excluding the extra bias unit which always outputs +1). As before, the training data will be loaded into the variables X and y. You have been provided with a set of network parameters

$(\Theta(1),\Theta(2))$ already trained by us. These are stored in ex3weights.mat and will be loaded by ex3 nn.m into

Theta1 and Theta2. The parameters have dimensions that are sized for a neural network with 25 units in the second layer and 10 output units (corresponding to the 10 digit classes)

%Load saved matrices
from   file
load('ex3weights.mat');

%The matrices Theta1 and Theta2 will now be in your Octave environment

%Theta1 has size 25 x 401 % Theta2 has size 26 x 10

$$\Theta^{(1)} \qquad \Theta^{(2)}$$

x_0

x_2

+1

+1

$h_\theta(x)$

$$a^{(1)} = x$$
$$(\text{add } a_0^{(1)})$$

$$z^{(2)} = \Theta^{(1)}a^{(1)}$$
$$a^{(2)} = g(z^{(2)})$$
$$(\text{add } a_0^{(2)})$$

$$z^{(3)} = \Theta^{(2)}a^{(2)}$$
$$a^{(3)} = g(z^{(3)}) = h_\theta(x)$$

Input Layer   Hidden Layer   Output Layer

2.2 Feedforward Propagation and Prediction

Now you will implement feedforward propagation for the neural network. You will need to complete the code in predict.m to return the neural network's prediction. You should implement the feedforward computation that computes $h_\theta(x(i))$ for every example i and returns the associated predictions. Similar to the one-vs-all classification strategy, the prediction from the neural network will be the label that has the largest output $(h_\theta(x))_k$.

Implementation Note: The matrix X contains the examples in rows. When you complete the code in predict.m, you will need to add the column of 1's to the matrix. The matrices Theta1 and Theta2 contain the parameters for each unit in rows. Specifically, the first row of Theta1 corresponds to the first hidden unit in the second layer. In Octave, when you compute $z^{(2)} = \Theta^{(1)} a^{(1)}$, be sure that you index (and if necessary, transpose) X correctly so that you get $a^{(l)}$ as a column vector.

Once you are done, ex3 nn.m will call your nnCostFunction using the loaded set of parameters for Theta1 and Theta2, and λ = 1. You should see that the accuracy is about 97.5%. After that, an interactive sequence will launch displaying images from the training set one at a time, while the console prints out the predicted label for the displayed image. To stop the image sequence, press Ctrl-C.

3. <u>Submission</u>

Please send us the final zip folder **<u>before Saturday, 25th June 2016 midnight</u>**. E-mails to be sent to both me and Akhilesh.